

▼ REQUIRED LIB

```

1 from tqdm import tqdm
2 import torch
3 import torch.nn as nn
4 import torch.nn.functional as F
5 import torch.optim as optim
6
7 import torchvision
8 from torchvision.datasets import MNIST
9 import torchvision.transforms as transforms
10 from torch.utils.data import DataLoader
11 import matplotlib.pyplot as plt
12
1
1 # Check if CUDA (GPU) is available, otherwise use CPU
2 device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
1
1 device
📄 device(type='cuda')

```

▼ Data Preprocessing and Loading Dataset

```

1 # transformations for the images
2 transform = transforms.ToTensor()
3
4 # Load the MNIST dataset and create data loaders
5 train_dataset = MNIST(root='./data', train=True, transform=transform, download=True)
6 test_dataset = MNIST(root='./data', train=False, transform=transform, download=True)
7
8 train_dataloader = DataLoader(dataset=train_dataset, batch_size=64, shuffle=True)
9 test_dataloader = DataLoader(dataset=test_dataset, batch_size=32, shuffle=False)

```

Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz>
 Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz> to ./data/MNIST/raw/train-images-idx3-ubyte.gz
 100%|██████████| 9912422/9912422 [00:00<00:00, 104129574.91it/s]
 Extracting ./data/MNIST/raw/train-images-idx3-ubyte.gz to ./data/MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz>
 Downloading <http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz> to ./data/MNIST/raw/train-labels-idx1-ubyte.gz
 100%|██████████| 28881/28881 [00:00<00:00, 41670345.31it/s]
 Extracting ./data/MNIST/raw/train-labels-idx1-ubyte.gz to ./data/MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz>
 Downloading <http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz> to ./data/MNIST/raw/t10k-images-idx3-ubyte.gz
 100%|██████████| 1648877/1648877 [00:00<00:00, 31281877.83it/s]
 Extracting ./data/MNIST/raw/t10k-images-idx3-ubyte.gz to ./data/MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz>
 Downloading <http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz> to ./data/MNIST/raw/t10k-labels-idx1-ubyte.gz
 100%|██████████| 4542/4542 [00:00<00:00, 14848424.60it/s]
 Extracting ./data/MNIST/raw/t10k-labels-idx1-ubyte.gz to ./data/MNIST/raw

```

1 train_dataset

```

Dataset MNIST
 Number of datapoints: 60000
 Root location: ./data
 Split: Train
 StandardTransform
 Transform: ToTensor()

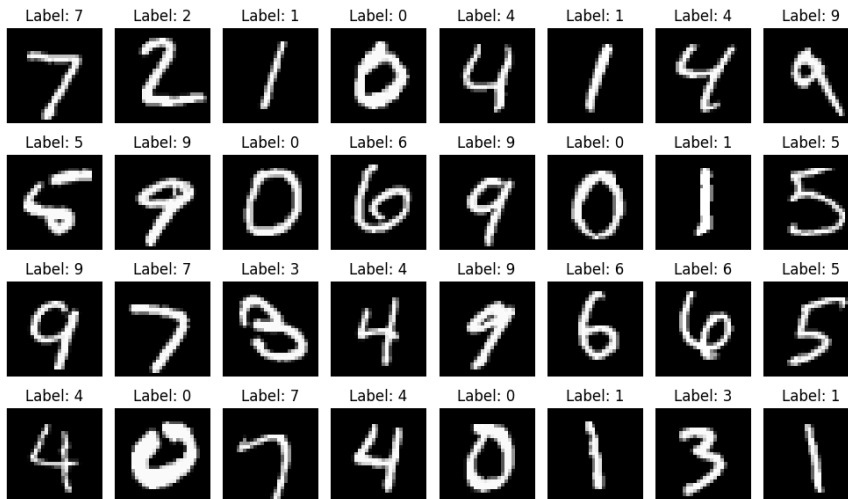
▼ Visualizing the Data

Let's visualize a few examples from the dataset to understand the input images better.

```

1 import matplotlib.pyplot as plt
2 import numpy as np
3
4 # Function to display a grid of images
5 def show_images(images, labels, nrow, ncol):
6     fig, axes = plt.subplots(nrow, ncol, figsize=(10, 6))
7     for i, ax in enumerate(axes.flat):
8         ax.imshow(np.squeeze(images[i]), cmap='gray')
9         ax.set_title(f"Label: {labels[i]}")
10        ax.axis('off')
11    plt.tight_layout()
12    plt.show()
13
14 # Get a batch of images and labels from the test dataset
15 images, labels = next(iter(test_dataloader))
16
17 # Display the images and their labels
18 show_images(images, labels, 4, 8)

```




```

1 img, label = train_dataset[0]
2 plt.imshow(np.squeeze(img), cmap='gray')

```

<matplotlib.image.AxesImage at 0x790162f065c0>



Model Architecture

Next, we'll define our CNN model architecture using PyTorch.

```

1
2 class DigitClassification(nn.Module):
3     def __init__(self):
4         super().__init__()
5         self.conv1 = nn.Conv2d(in_channels=1, out_channels=16, kernel_size=5)
6         self.pool = nn.MaxPool2d(2)
7         self.conv2 = nn.Conv2d(in_channels=16, out_channels=8, kernel_size=3)
8         self.L1 = nn.Linear(in_features=8*5*5, out_features=64)
9         self.L2 = nn.Linear(in_features=64, out_features=10)
10
11     def forward(self, x):
12         x = self.conv1(x) # 1*28*28 --> 16*24*24
13         x = self.pool(x)
14         x = F.relu(x)
15
16         x = self.conv2(x)
17         x = self.pool(x)
18         x = F.relu(x)
19
20         x = x.flatten(start_dim=1, end_dim=-1)
21         x = F.relu(self.L1(x))
22         x = F.relu(self.L2(x))
23         return x

```

✓ Training the Model

Define Loss Function and Optimizer We need to specify the loss function and optimizer for training the model.

```

1 # Instantiate the model and move it to the appropriate device
2 model = DigitClassification().to(device)
3
4 # Define the loss function and optimizer
5 criterion = nn.CrossEntropyLoss()
6 optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9)

```

✓ Training Loop

Now, we'll implement the training loop to train our model on the MNIST dataset.

```

1 # Define the training function
2 def train(model, train_dataloader, criterion, optimizer):
3     total_loss = 0
4     for batch in train_dataloader:
5         data, label = batch[0].to(device), batch[1].to(device) # Move data to device
6
7         optimizer.zero_grad()
8
9         predicted = model(data)
10        loss = criterion(predicted, label)
11
12        loss.backward()
13        optimizer.step()
14        total_loss += loss.item()
15    return total_loss
16

```

✓ TEST loop

```

1 def test(model, test_dataloader, criterion, optimizer):
2     total_loss = 0
3     with torch.no_grad(): # No need to track gradients during testing
4         for batch in test_dataloader:
5             data, label = batch[0].to(device), batch[1].to(device) # Move data to device
6
7             predicted = model(data)
8             loss = criterion(predicted, label)
9
10            total_loss += loss.item()
11    return total_loss

```

```

1 # Train and test the model for a certain number of epochs
2 num_epochs = 10
3 training_losses = []
4 testing_losses = []

1 for epoch in tqdm(range(num_epochs), desc='Training Progress'):
2
3     training_loss = train(model, train_dataloader, criterion, optimizer)
4     testing_loss = test(model, test_dataloader, criterion, optimizer)
5
6     # Append losses for visualization
7     training_losses.append(training_loss)
8     testing_losses.append(testing_loss)
9
10    print(f"Training Loss: {training_loss}, Testing Loss: {testing_loss}")


Training Progress: 10%|██████| 1/10 [00:11<01:41, 11.26s/it]Training Loss: 1833.5078232884407, Testing Loss: 305.89
Training Progress: 20%|██████| 2/10 [00:20<01:20, 10.04s/it]Training Loss: 448.25220078229904, Testing Loss: 92.100
Training Progress: 30%|██████| 3/10 [00:29<01:07, 9.60s/it]Training Loss: 250.6873406395316, Testing Loss: 65.9009
Training Progress: 40%|██████| 4/10 [00:39<00:58, 9.72s/it]Training Loss: 181.1206422690302, Testing Loss: 46.1270
Training Progress: 50%|██████| 5/10 [00:48<00:46, 9.38s/it]Training Loss: 140.67512695118785, Testing Loss: 39.108
Training Progress: 60%|██████| 6/10 [00:57<00:37, 9.28s/it]Training Loss: 115.48206906672567, Testing Loss: 32.162
Training Progress: 70%|██████| 7/10 [01:06<00:27, 9.23s/it]Training Loss: 99.26806882582605, Testing Loss: 28.6702
Training Progress: 80%|██████| 8/10 [01:15<00:18, 9.21s/it]Training Loss: 87.5746425203979, Testing Loss: 26.31545
Training Progress: 90%|██████| 9/10 [01:24<00:08, 8.97s/it]Training Loss: 79.41148270107806, Testing Loss: 22.1120
Training Progress: 100%|██████| 10/10 [01:33<00:00, 9.31s/it]Training Loss: 72.75542741268873, Testing Loss: 21.402

```

```

1 # Plotting training and testing losses
2 plt.plot(training_losses, label='Training Loss')
3 plt.plot(testing_losses, label='Testing Loss')
4 plt.xlabel('Epoch')
5 plt.ylabel('Loss')
6 plt.title('Training and Testing Loss')
7 plt.legend()
8 plt.show()

```



✓ Predictions and Visualization

Making Predictions

Let's make predictions on new data using the trained model.

```

1 # Function to make predictions on new data
2 def predict(model, dataloader):
3     predictions = []
4     with torch.no_grad():
5         for data in dataloader:
6             inputs = data[0].to(device)
7             outputs = model(inputs)
8             _, predicted = torch.max(outputs.data, 1)
9             predictions.extend(predicted.cpu().numpy())
10    return predictions
11

```

```

1 # Make predictions on the test dataset
2 test_predictions = predict(model, test_dataloader)

```

Visualizing Predictions

We'll visualize a random sample of test images along with their predicted labels.

```

1 # Function to plot a random sample of images along with their predicted labels
2 def visualize_predictions(dataset, predictions, num_samples=5):
3     plt.figure(figsize=(15, 7))
4     samples = np.random.choice(len(dataset), num_samples, replace=False)
5     for i, idx in enumerate(samples):
6         plt.subplot(1, num_samples, i + 1)
7         image, label = dataset[idx]
8         plt.imshow(image.squeeze(), cmap='gray')
9         plt.title(f"Predicted: {predictions[idx]}, Actual: {label}")
10        plt.axis('off')
11    plt.show()
12

```

```

1 # Visualize predictions on a random sample of test images
2 visualize_predictions(test_dataset, test_predictions)

```

